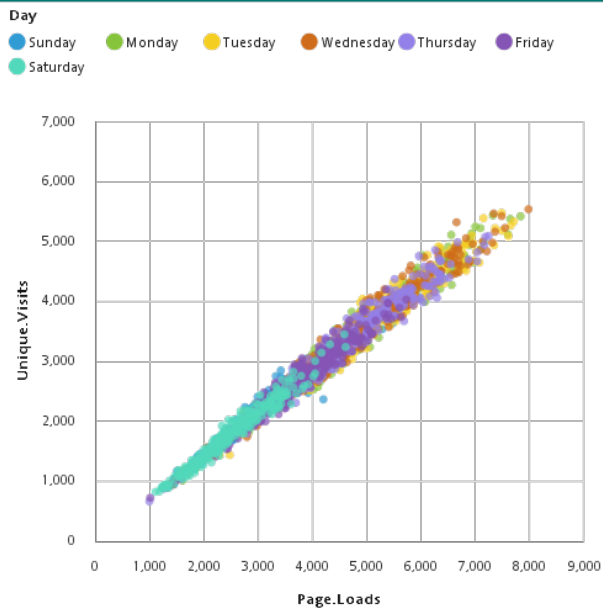



page loads in each day of week




Day

- Sunday
- Monday
- Tuesday
- Wednesday
- Thursday
- Friday
- Saturday




unique visits




visualizing the first time visits in overall page loads

First.Time.Visits

456,449 907,752





website-visitors-analysis

November 5, 2023

```
[3]: import pandas as pd
file_loc = '/daily-website-visitors.csv'
dataset = pd.read_csv(file_loc,
                      index_col = 'Date',
                      thousands=',')
dataset.index = pd.to_datetime(dataset.index)
dataset
```

```
[3]:
```

	Row	Day	Day.Of.Week	Page.Loads	Unique.Visits	\
Date						
2014-09-14	1	Sunday	1	2146	1582	
2014-09-15	2	Monday	2	3621	2528	
2014-09-16	3	Tuesday	3	3698	2630	
2014-09-17	4	Wednesday	4	3667	2614	
2014-09-18	5	Thursday	5	3316	2366	
...	
2020-08-15	2163	Saturday	7	2221	1696	
2020-08-16	2164	Sunday	1	2724	2037	
2020-08-17	2165	Monday	2	3456	2638	
2020-08-18	2166	Tuesday	3	3581	2683	
2020-08-19	2167	Wednesday	4	2064	1564	

	First.Time.Visits	Returning.Visits
Date		
2014-09-14	1430	152
2014-09-15	2297	231
2014-09-16	2352	278
2014-09-17	2327	287
2014-09-18	2130	236
...
2020-08-15	1373	323
2020-08-16	1686	351
2020-08-17	2181	457
2020-08-18	2184	499
2020-08-19	1297	267

[2167 rows x 7 columns]

```
[4]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 2167 entries, 2014-09-14 to 2020-08-19
Data columns (total 7 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Row                    2167 non-null  int64
 1   Day                    2167 non-null  object
 2   Day.Of.Week            2167 non-null  int64
 3   Page.Loads             2167 non-null  int64
 4   Unique.Visits          2167 non-null  int64
 5   First.Time.Visits      2167 non-null  int64
 6   Returning.Visits       2167 non-null  int64
dtypes: int64(6), object(1)
memory usage: 135.4+ KB
```

```
[5]: dataset.describe()
```

```
[5]:
```

	Row	Day.Of.Week	Page.Loads	Unique.Visits	\
count	2167.000000	2167.000000	2167.000000	2167.000000	
mean	1084.000000	3.997231	4116.989386	2943.646516	
std	625.703338	2.000229	1350.977843	977.886472	
min	1.000000	1.000000	1002.000000	667.000000	
25%	542.500000	2.000000	3114.500000	2226.000000	
50%	1084.000000	4.000000	4106.000000	2914.000000	
75%	1625.500000	6.000000	5020.500000	3667.500000	
max	2167.000000	7.000000	7984.000000	5541.000000	

	First.Time.Visits	Returning.Visits
count	2167.000000	2167.000000
mean	2431.824181	511.822335
std	828.704688	168.736370
min	522.000000	133.000000
25%	1830.000000	388.500000
50%	2400.000000	509.000000
75%	3038.000000	626.500000
max	4616.000000	1036.000000

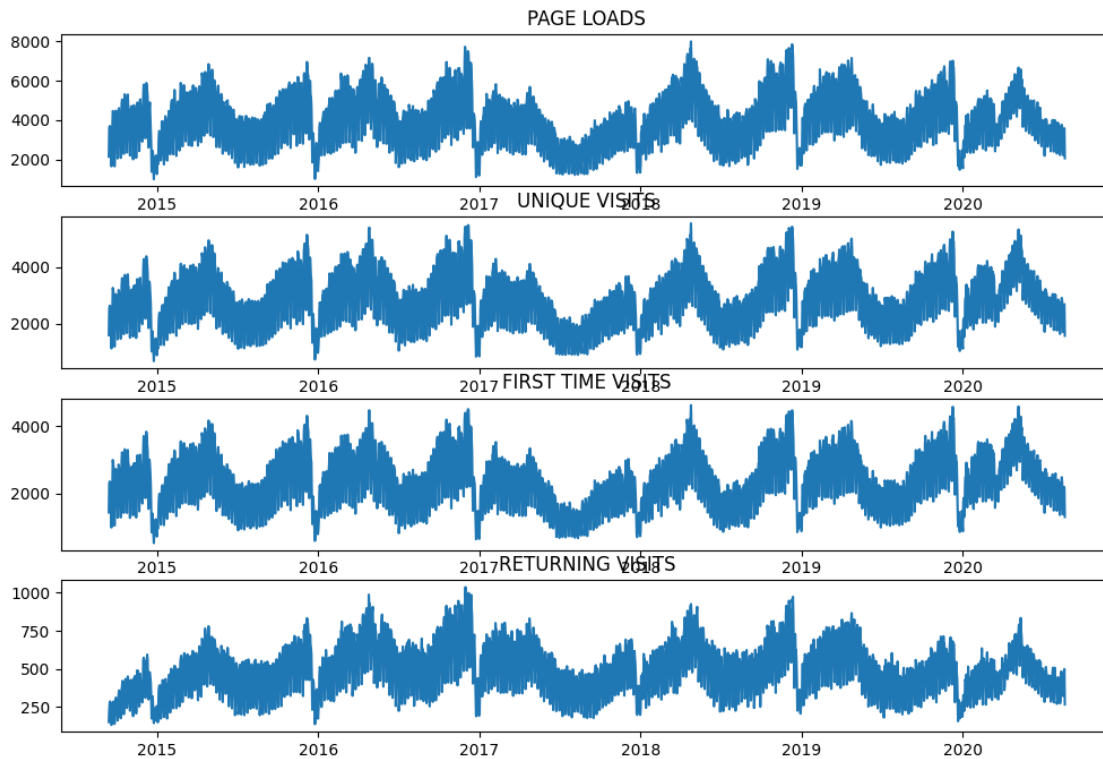
#LET'S make a little visual about the visits.

```
[13]: import matplotlib.pyplot as plt

fig, ax = plt.subplots(4, figsize=(12, 8))
ax[0].plot(dataset['Page.Loads'])
ax[1].plot(dataset['Unique.Visits'])
ax[2].plot(dataset['First.Time.Visits'])
ax[3].plot(dataset['Returning.Visits'])
```

```
ax[0].set_title("PAGE LOADS")
ax[1].set_title("UNIQUE VISITS")
ax[2].set_title("FIRST TIME VISITS")
ax[3].set_title("RETURNING VISITS")

plt.show()
```




#LET'S PREPROCESS THE DATA.#

```
[57]: target_col = dataset['Returning.Visits']
      target_col
```

```
[57]: Date
      2014-09-14    152
      2014-09-15    231
      2014-09-16    278
      2014-09-17    287
      2014-09-18    236
      ...
      2020-08-15    323
      2020-08-16    351
      2020-08-17    457
      2020-08-18    499
```

2020-08-19 267
Name: Returning.Visits, Length: 2167, dtype: int64

```
[23]: target_col.plot(figsize=(15, 3))  
plt.show()
```



```
[26]: len(target_col)
```

[26]: 2167

```
[25]: TEST_DATA_PERCENTAGE = 0.1  
boundary_indx = int((1 - TEST_DATA_PERCENTAGE) * len(target_col))  
  
train_data = target_col[:boundary_indx]  
test_data = target_col[boundary_indx:]  
  
print(f"Train data:\tReturning Visits [{boundary_indx}] ({boundary_indx})")  
print(f"Test data:\tReturning Visits [{boundary_indx}:] ({len(target_col) -  
↪ boundary_indx})")  
print(f"\nLast target on train data: {train_data[-1]}")
```

Train data: Returning Visits [:1950] (1950)
Test data: Returning Visits [1950:] (217)

Last target on train data: 394

```
[33]: print(f"Train data ending values:{target_col[TEST_DATA_BOUNDARY_INDEX - 10:↪  
↪ TEST_DATA_BOUNDARY_INDEX].values}")  
print(f"Test dataset starting valuesL {target_col[TEST_DATA_BOUNDARY_INDEX:↪  
↪ TEST_DATA_BOUNDARY_INDEX + 10].values}")
```

Train data ending values:[429 423 442 464 372 253 277 515 434 394]
Test dataset starting valuesL [441 413 246 314 443 484 473 490 353 249]

#WINDOW_IZE THE DATASET

```
[58]: import numpy as np
import pandas as pd

if isinstance(target_col, pd.DataFrame) or isinstance(target_col, pd.Series):
    target_array = target_col.values
else:
    target_array = np.array(target_col)
WINDOW_SIZE = 3
# Now, you can use target_array in your code
train_dataset = timeseries_dataset_from_array(target_array[: -WINDOW_SIZE],
                                              target_array[WINDOW_SIZE:],
                                              sequence_length=WINDOW_SIZE,

                                              ↪end_index=TEST_DATA_BOUNDARY_INDEX - 1)
```

```
[60]: from keras.utils import timeseries_dataset_from_array
import numpy as np

WINDOW_SIZE = 3

# Ensure target_col is a NumPy array
target_col1 = np.array(target_col)

# Create the train_set
train_set = timeseries_dataset_from_array(
    data=target_col1[: -WINDOW_SIZE],
    targets=target_col1[WINDOW_SIZE:],
    sequence_length=WINDOW_SIZE,
    end_index=TEST_DATA_BOUNDARY_INDEX - WINDOW_SIZE)
```

```
[59]: target_col[TEST_DATA_BOUNDARY_INDEX-10:TEST_DATA_BOUNDARY_INDEX+10].values, ↪
↪(list(train_dataset)[-1][0][-1].numpy(), list(train_dataset)[-1][1][-1].
↪numpy())
```

```
[59]: (array([429, 423, 442, 464, 372, 253, 277, 515, 434, 394, 441, 413, 246,
314, 443, 484, 473, 490, 353, 249]),
(array([277, 515, 434]), 394))
```

```
[62]: test_dataset = ↪
↪timeseries_dataset_from_array(target_col[TEST_DATA_BOUNDARY_INDEX - ↪
↪WINDOW_SIZE:],
↪
↪target_col[TEST_DATA_BOUNDARY_INDEX:],
↪
↪sequence_length=WINDOW_SIZE
)
len(test_dataset), len(list(test_dataset.unbatch()))
(2, 217)
```

```
[62]: (2, 217)
```

```
[64]: target_col[TEST_DATA_BOUNDARY_INDEX-10:TEST_DATA_BOUNDARY_INDEX+10].values,   
      ↪list(test_dataset)[0][0][0].numpy(), list(test_dataset)[0][1][0].numpy()
```

```
[64]: (array([429, 423, 442, 464, 372, 253, 277, 515, 434, 394, 441, 413, 246,   
            314, 443, 484, 473, 490, 353, 249]),   
      array([515, 434, 394]),   
      441)
```

first point in test dataset

```
[65]: list(test_dataset)[0][0][0].numpy(), list(test_dataset)[0][1][0].numpy()
```

```
[65]: (array([515, 434, 394]), 441)
```

last point in test dataset

```
[66]: list(test_dataset)[-1][0][-1].numpy(), list(test_dataset)[-1][1][-1].numpy()
```

```
[66]: (array([351, 457, 499]), 267)
```


LET'S PLOT THE TRAIN AND TEST DATA

```
[68]: import matplotlib.dates as mdates

def plot_time_series(predictions = None, start_index=1500):
    timesteps = pd.to_datetime(target_col.index)

    fig, ax = plt.subplots(1, figsize=(15, 5))
    ax.xaxis.set_major_locator(mdates.MonthLocator(bymonth=(1, 7)))
    ax.xaxis.set_minor_locator(mdates.MonthLocator())
    ax.xaxis.set_major_formatter(mdates.DateFormatter('%Y-%b'))
    #plot train data
    plt.plot(timesteps[start_index:TEST_DATA_BOUNDARY_INDEX],   
      ↪target_col[start_index:TEST_DATA_BOUNDARY_INDEX],
            color='blue')
    #plot test dataset
    plt.plot(timesteps[TEST_DATA_BOUNDARY_INDEX:],   
      ↪target_col[TEST_DATA_BOUNDARY_INDEX:],
            color='green', linewidth=0.4)
    if predictions is not None:
        pred_timesteps = timesteps[TEST_DATA_BOUNDARY_INDEX:]
        plt.plot(pred_timesteps, predictions, linewidth=0.4, color='red')
        plt.scatter(pred_timesteps, predictions, s=0.4, color='red')

plot_time_series()
```

#BASELINE MODEL

```
[70]: import tensorflow as tf
from keras.layers import Layer
from keras import Model

class NaiveForecastLayer(Model):
    def __init__(self):
        super().__init__()

    def call(self, inputs):
        result = inputs[:, -1]
        return result[:, tf.newaxis]
```


```
[71]: baseline_model = NaiveForecastLayer()
baseline_model._name = 'model_0'

baseline_model.compile(metrics=[tf.keras.metrics.MeanAbsoluteError()])
```

```
[72]: baseline_predictions = baseline_model.predict(test_dataset)
```

2/2 [=====] - 0s 20ms/step

```
[75]: plot_time_series(baseline_predictions.ravel(), start_index=1900)
```



```
[77]: y_true = target_col[TEST_DATA_BOUNDARY_INDEX : ]

len(y_true), y_true
```

```
[77]: (217,
      Date
      2020-01-16    441
      2020-01-17    413
      2020-01-18    246
      2020-01-19    314
      2020-01-20    443
      ...
      2020-08-15    323
      2020-08-16    351
      2020-08-17    457
      2020-08-18    499
      2020-08-19    267
      Name: Returning.Visits, Length: 217, dtype: int64)
```

```
[78]: from sklearn.metrics import mean_absolute_error, mean_squared_error, \
      ↪ mean_absolute_percentage_error

def evaluate_predictions(y_true, y_preds):
    mae = mean_absolute_error(y_true, y_preds)
    mse = mean_squared_error(y_true, y_preds)
    rmse = np.sqrt(mse)
    mape = mean_absolute_percentage_error(y_true, y_preds)

    return {
        'mae': mae,
        'mse': mse,
        'rmse': rmse,
        'mape': mape
```

```
}
```

```
evaluate_predictions(y_true, baseline_predictions)
```


```
[78]: {'mae': 72.19815668202764,  
      'mse': 8508.622119815669,  
      'rmse': 92.24219273096054,  
      'mape': 0.16713927858326993}
```

```
[80]: MODEL_METRICS = pd.DataFrame(columns=['mae', 'mse', 'rmse', 'mape'])
```

```
def evaluate_model(model):  
    predictions = model.predict(test_dataset, verbose=0)  
    metrics = evaluate_predictions(y_true, predictions)  
  
    MODEL_METRICS.loc[model.name] = metrics  
    plot_time_series(predictions.ravel(), start_index=1900)  
    return metrics
```

```
[81]: evaluate_model(baseline_model)
```

```
[81]: {'mae': 72.19815668202764,  
      'mse': 8508.622119815669,  
      'rmse': 92.24219273096054,  
      'mape': 0.16713927858326993}
```



```
[82]: MODEL_METRICS
```

```
[82]:
```

	mae	mse	rmse	mape
model_0	72.198157	8508.62212	92.242193	0.167139

#Recurrent Network Model

```
[83]: from keras.layers import GRU, Dense, Input, Lambda
      from keras import Sequential
```

```
tf.random.set_seed(42)
model_1 = Sequential([
    Input(shape=(WINDOW_SIZE,)),
    Lambda(lambda x: tf.expand_dims(x, axis=1)),
    GRU(128, activation="relu"),
    Dense(1)
], name='model_1')

model_1.compile(
    loss=tf.keras.losses.MeanAbsoluteError(),
    optimizer=tf.keras.optimizers.Adam()
)

model_1.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #
lambda (Lambda)	(None, 1, 3)	0
gru (GRU)	(None, 128)	51072
dense (Dense)	(None, 1)	129

=====
Total params: 51201 (200.00 KB)
Trainable params: 51201 (200.00 KB)
Non-trainable params: 0 (0.00 Byte)
=====

```
[84]: from keras.callbacks import ModelCheckpoint
      import os

      def create_checkpoint_callback(model):
          filepath = os.path.join('models', model.name)
          return ModelCheckpoint(filepath, monitor='loss', save_weights_only=True,
                                  ↪save_best_only=True)

      model_1.fit(train_dataset, epochs=5, callbacks=[
          ↪create_checkpoint_callback(model_1) ])
```

Epoch 1/5


16/16 [=====] - 3s 32ms/step - loss: 478.9003

```
Epoch 2/5
16/16 [=====] - 1s 31ms/step - loss: 302.9468
Epoch 3/5
16/16 [=====] - 0s 28ms/step - loss: 170.8370
Epoch 4/5
16/16 [=====] - 0s 28ms/step - loss: 133.6493
Epoch 5/5
16/16 [=====] - 0s 22ms/step - loss: 130.8549
```

[84]: <keras.src.callbacks.History at 0x7be4abaff850>

```
[85]: evaluate_model(model_1)
```

```
[85]: {'mae': 88.47757000198013,
      'mse': 11614.890892560847,
      'rmse': 107.77240320490607,
      'mape': 0.1947460150386166}
```



```
[86]: MODEL_METRICS
```

```
[86]:
```

	mae	mse	rmse	mape
model_0	72.198157	8508.622120	92.242193	0.167139
model_1	88.477570	11614.890893	107.772403	0.194746

```
[87]: unbatched_train_dataset = dataset[:,TEST_DATA_BOUNDARY_INDEX + 1].copy()
      unbatched_train_dataset
```

```
[87]:
```

	Row	Day	Day.Of.Week	Page.Loads	Unique.Visits	\
Date						
2014-09-14	1	Sunday	1	2146	1582	
2014-09-15	2	Monday	2	3621	2528	
2014-09-16	3	Tuesday	3	3698	2630	
2014-09-17	4	Wednesday	4	3667	2614	

2014-09-18	5	Thursday	5	3316	2366
...
2020-01-12	1947	Sunday	1	2762	2238
2020-01-13	1948	Monday	2	4298	3242
2020-01-14	1949	Tuesday	3	3838	2884
2020-01-15	1950	Wednesday	4	3754	2864
2020-01-16	1951	Thursday	5	3817	2951

Date	First.Time.Visits	Returning.Visits
2014-09-14	1430	152
2014-09-15	2297	231
2014-09-16	2352	278
2014-09-17	2327	287
2014-09-18	2130	236
...
2020-01-12	1961	277
2020-01-13	2727	515
2020-01-14	2450	434
2020-01-15	2470	394
2020-01-16	2510	441


[1951 rows x 7 columns]

```
[88]: dataset_by_day = unbatched_train_dataset.groupby(by=['Day'])
dataset_by_day['Returning.Visits'].mean()
```

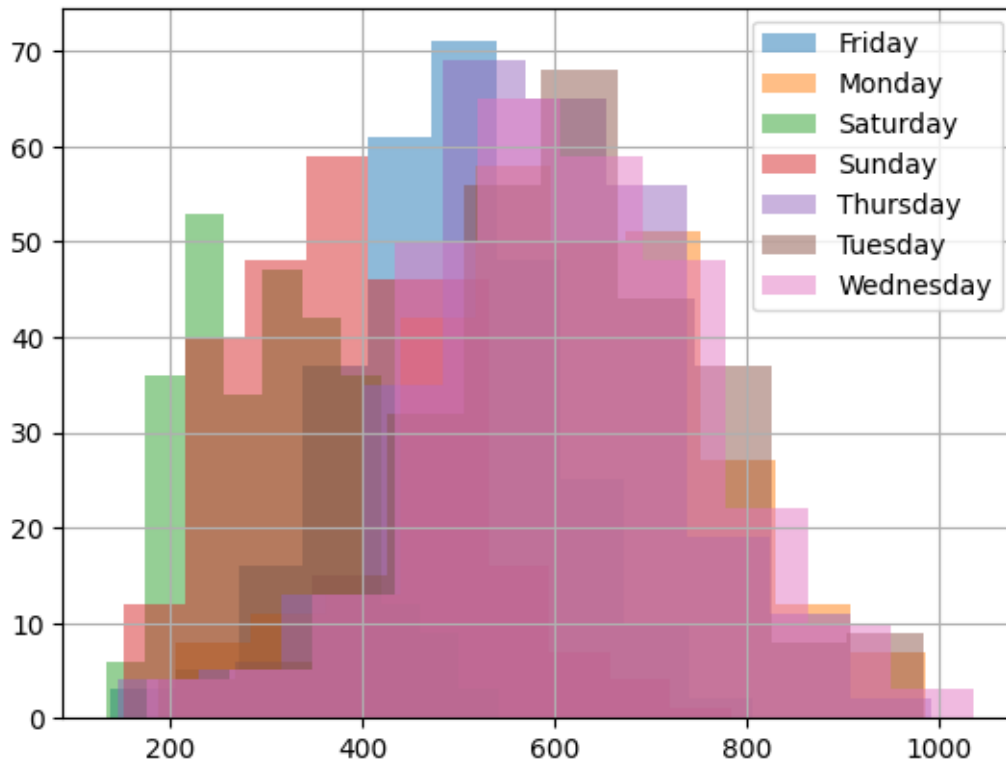
```
[88]: Day
Friday      484.697842
Monday      606.512545
Saturday    306.071942
Sunday      390.573477
Thursday    584.627240
Tuesday     617.888889
Wednesday   614.369176
Name: Returning.Visits, dtype: float64
```

```
[89]: DAYS_OF_WEEK = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
↳ 'Saturday', 'Sunday']
pd.DataFrame(dataset_by_day['Returning.Visits'].mean()).loc[DAYS_OF_WEEK].
↳ plot(kind='bar')
```

```
[89]: <Axes: xlabel='Day'>
```



```
[90]: dataset_by_day['Returning.Visits'].hist(legend=True, alpha=0.5)  
plt.show()
```



```
[91]: import calendar
```

```
train_dataset_with_months = unbatched_train_dataset.copy()
train_dataset_with_months['Month.Name'] = pd.Series(train_dataset_with_months.
    ↪ index,
                                                    ↪
    ↪ index=train_dataset_with_months.index)\
                                                    .apply(lambda x: calendar.
    ↪ month_name[x.month])
train_dataset_with_months
```

```
[91]:
```

	Row	Day	Day.Of.Week	Page.Loads	Unique.Visits	\
Date						
2014-09-14	1	Sunday	1	2146	1582	
2014-09-15	2	Monday	2	3621	2528	
2014-09-16	3	Tuesday	3	3698	2630	
2014-09-17	4	Wednesday	4	3667	2614	
2014-09-18	5	Thursday	5	3316	2366	
...	
2020-01-12	1947	Sunday	1	2762	2238	
2020-01-13	1948	Monday	2	4298	3242	
2020-01-14	1949	Tuesday	3	3838	2884	

2020-01-15	1950	Wednesday	4	3754	2864
2020-01-16	1951	Thursday	5	3817	2951

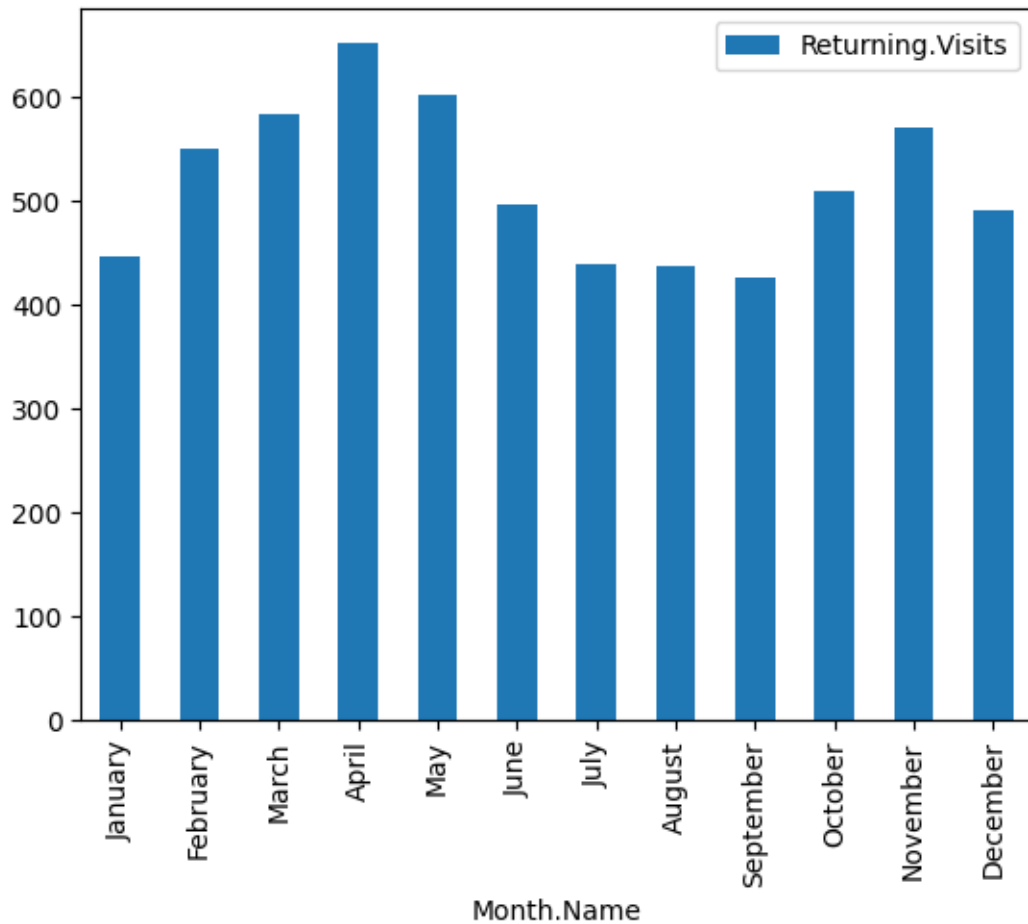
Date	First.Time.Visits	Returning.Visits	Month.Name
2014-09-14	1430	152	September
2014-09-15	2297	231	September
2014-09-16	2352	278	September
2014-09-17	2327	287	September
2014-09-18	2130	236	September
...
2020-01-12	1961	277	January
2020-01-13	2727	515	January
2020-01-14	2450	434	January
2020-01-15	2470	394	January
2020-01-16	2510	441	January

[1951 rows x 8 columns]

```
[92]: MONTH_NAMES = list(calendar.month_name)[1:]
dataset_group_by_month = train_dataset_with_months.groupby(by='Month.Name')
dataset_group_by_month['Returning.Visits'].mean().loc[MONTH_NAMES]
```

```
[92]: Month.Name
January      445.976608
February     549.354610
March        583.470968
April        651.740000
May          601.135484
June         496.180000
July         438.509677
August       437.522581
September    426.173653
October      509.209677
November     569.716667
December     490.274194
Name: Returning.Visits, dtype: float64
```

```
[93]: pd.DataFrame(dataset_group_by_month['Returning.Visits'].mean()).
      ↪loc[MONTH_NAMES].plot(kind='bar')
plt.show()
```



Prepare the dataset

```
[94]: train_dataset_with_months
```

```
[94]:
```

	Row	Day	Day.Of.Week	Page.Loads	Unique.Visits	\
Date						
2014-09-14	1	Sunday	1	2146	1582	
2014-09-15	2	Monday	2	3621	2528	
2014-09-16	3	Tuesday	3	3698	2630	
2014-09-17	4	Wednesday	4	3667	2614	
2014-09-18	5	Thursday	5	3316	2366	
...	
2020-01-12	1947	Sunday	1	2762	2238	
2020-01-13	1948	Monday	2	4298	3242	
2020-01-14	1949	Tuesday	3	3838	2884	
2020-01-15	1950	Wednesday	4	3754	2864	
2020-01-16	1951	Thursday	5	3817	2951	

	First.Time.Visits	Returning.Visits	Month.Name
Date			
2014-09-14	1430	152	September
2014-09-15	2297	231	September
2014-09-16	2352	278	September
2014-09-17	2327	287	September
2014-09-18	2130	236	September
...
2020-01-12	1961	277	January
2020-01-13	2727	515	January
2020-01-14	2450	434	January
2020-01-15	2470	394	January
2020-01-16	2510	441	January

[1951 rows x 8 columns]

```
[95]: dataset2 = train_dataset_with_months.copy()[['Day', 'Month.Name', 'Returning.
↳Visits']]
dataset2
```

```
[95]:
```

	Day	Month.Name	Returning.Visits
Date			
2014-09-14	Sunday	September	152
2014-09-15	Monday	September	231
2014-09-16	Tuesday	September	278
2014-09-17	Wednesday	September	287
2014-09-18	Thursday	September	236
...
2020-01-12	Sunday	January	277
2020-01-13	Monday	January	515
2020-01-14	Tuesday	January	434
2020-01-15	Wednesday	January	394
2020-01-16	Thursday	January	441

[1951 rows x 3 columns]

```
[96]: def windowize_dataset(dataset):
        for i in range(WINDOW_SIZE):
            dataset[f'Returning.Visits[t-{i+1}]]' = dataset['Returning.Visits'].
↳shift(periods=i+1)
        return dataset

dataset2 = windowize_dataset(dataset2.copy())
dataset2
```

```
[96]:
```

	Day	Month.Name	Returning.Visits	Returning.Visits[t-1]	\
Date					

2014-09-14	Sunday	September	152	NaN
2014-09-15	Monday	September	231	152.0
2014-09-16	Tuesday	September	278	231.0
2014-09-17	Wednesday	September	287	278.0
2014-09-18	Thursday	September	236	287.0
...
2020-01-12	Sunday	January	277	253.0
2020-01-13	Monday	January	515	277.0
2020-01-14	Tuesday	January	434	515.0
2020-01-15	Wednesday	January	394	434.0
2020-01-16	Thursday	January	441	394.0

	Returning.Visits[t-2]	Returning.Visits[t-3]
Date		
2014-09-14	NaN	NaN
2014-09-15	NaN	NaN
2014-09-16	152.0	NaN
2014-09-17	231.0	152.0
2014-09-18	278.0	231.0
...
2020-01-12	372.0	464.0
2020-01-13	253.0	372.0
2020-01-14	277.0	253.0
2020-01-15	515.0	277.0
2020-01-16	434.0	515.0

[1951 rows x 6 columns]

```
[97]: dataset2 = dataset2.dropna()
dataset2
```

```
[97]:
```

	Day	Month.Name	Returning.Visits	Returning.Visits[t-1]	\
Date					
2014-09-17	Wednesday	September	287	278.0	
2014-09-18	Thursday	September	236	287.0	
2014-09-19	Friday	September	241	236.0	
2014-09-20	Saturday	September	133	241.0	
2014-09-21	Sunday	September	175	133.0	
...	
2020-01-12	Sunday	January	277	253.0	
2020-01-13	Monday	January	515	277.0	
2020-01-14	Tuesday	January	434	515.0	
2020-01-15	Wednesday	January	394	434.0	
2020-01-16	Thursday	January	441	394.0	

	Returning.Visits[t-2]	Returning.Visits[t-3]
Date		

2014-09-17	231.0	152.0
2014-09-18	278.0	231.0
2014-09-19	287.0	278.0
2014-09-20	236.0	287.0
2014-09-21	241.0	236.0
...
2020-01-12	372.0	464.0
2020-01-13	253.0	372.0
2020-01-14	277.0	253.0
2020-01-15	515.0	277.0
2020-01-16	434.0	515.0

[1948 rows x 6 columns]

```
[98]: rv_cols = [f"Returning.Visits[t-{i+1}]" for i in range(WINDOW_SIZE)]
```

```
dataset2_rv_history_features = dataset2[rv_cols]
dataset2_rv_history_features
```

```
[98]:          Returning.Visits[t-1]  Returning.Visits[t-2]  \
```

Date		
2014-09-17	278.0	231.0
2014-09-18	287.0	278.0
2014-09-19	236.0	287.0
2014-09-20	241.0	236.0
2014-09-21	133.0	241.0
...
2020-01-12	253.0	372.0
2020-01-13	277.0	253.0
2020-01-14	515.0	277.0
2020-01-15	434.0	515.0
2020-01-16	394.0	434.0

```
          Returning.Visits[t-3]
```

Date	
2014-09-17	152.0
2014-09-18	231.0
2014-09-19	278.0
2014-09-20	287.0
2014-09-21	236.0
...	...
2020-01-12	464.0
2020-01-13	372.0
2020-01-14	253.0
2020-01-15	277.0
2020-01-16	515.0

[1948 rows x 3 columns]

```
[99]: dataset2_cat_features = dataset2[['Day', 'Month.Name']]
dataset2_cat_features
```

```
[99]:
```

	Day	Month.Name
Date		
2014-09-17	Wednesday	September
2014-09-18	Thursday	September
2014-09-19	Friday	September
2014-09-20	Saturday	September
2014-09-21	Sunday	September
...
2020-01-12	Sunday	January
2020-01-13	Monday	January
2020-01-14	Tuesday	January
2020-01-15	Wednesday	January
2020-01-16	Thursday	January

[1948 rows x 2 columns]

```
[100]: train_dataset2 = dataset2['Returning.Visits']
train_dataset2
```

```
[100]:
```

Date	
2014-09-17	287
2014-09-18	236
2014-09-19	241
2014-09-20	133
2014-09-21	175
...	
2020-01-12	277
2020-01-13	515
2020-01-14	434
2020-01-15	394
2020-01-16	441

Name: Returning.Visits, Length: 1948, dtype: int64

Building the model

```
[101]: from keras.layers import Concatenate, Dropout

tf.random.set_seed(42)
def build_model_3():
    seq_input = Input(shape=(WINDOW_SIZE,))
    lambda_layer = Lambda(lambda x: x[:, tf.newaxis])(seq_input)
    rnn_layer = GRU(64, activation='relu')(lambda_layer)
```

```

cat_input = Input(shape=(2,))
cat_dense_layer = Dense(32, activation='relu')(cat_input)

concat_layer = Concatenate()([rnn_layer, cat_dense_layer])
dense_layer1 = Dense(128, activation='relu')(concat_layer)
dropout_layer = Dropout(0.5)(dense_layer1)
output_layer = Dense(1, activation='linear')(dropout_layer)

return Model(inputs=[seq_input, cat_input], outputs=output_layer,
↳name="model_3")

model_3 = build_model_3()
model_3.compile(
    loss=tf.keras.losses.MeanAbsoluteError(),
    optimizer=tf.keras.optimizers.Adam()
)

model_3.summary()

```

Model: "model_3"

Layer (type)	Output Shape	Param #	Connected to
input_2 (InputLayer)	[(None, 3)]	0	[]
lambda_1 (Lambda) ['input_2[0][0]']	(None, 1, 3)	0	
input_3 (InputLayer)	[(None, 2)]	0	[]
gru_1 (GRU) ['lambda_1[0][0]']	(None, 64)	13248	
dense_1 (Dense) ['input_3[0][0]']	(None, 32)	96	
concatenate (Concatenate) ['gru_1[0][0]', 'dense_1[0][0]']	(None, 96)	0	
dense_2 (Dense) ['concatenate[0][0]']	(None, 128)	12416	
dropout (Dropout)	(None, 128)	0	

```
['dense_2[0][0]']
```

```
dense_3 (Dense)          (None, 1)          129  
['dropout[0][0]']
```

```
=====
```

```
=====
```

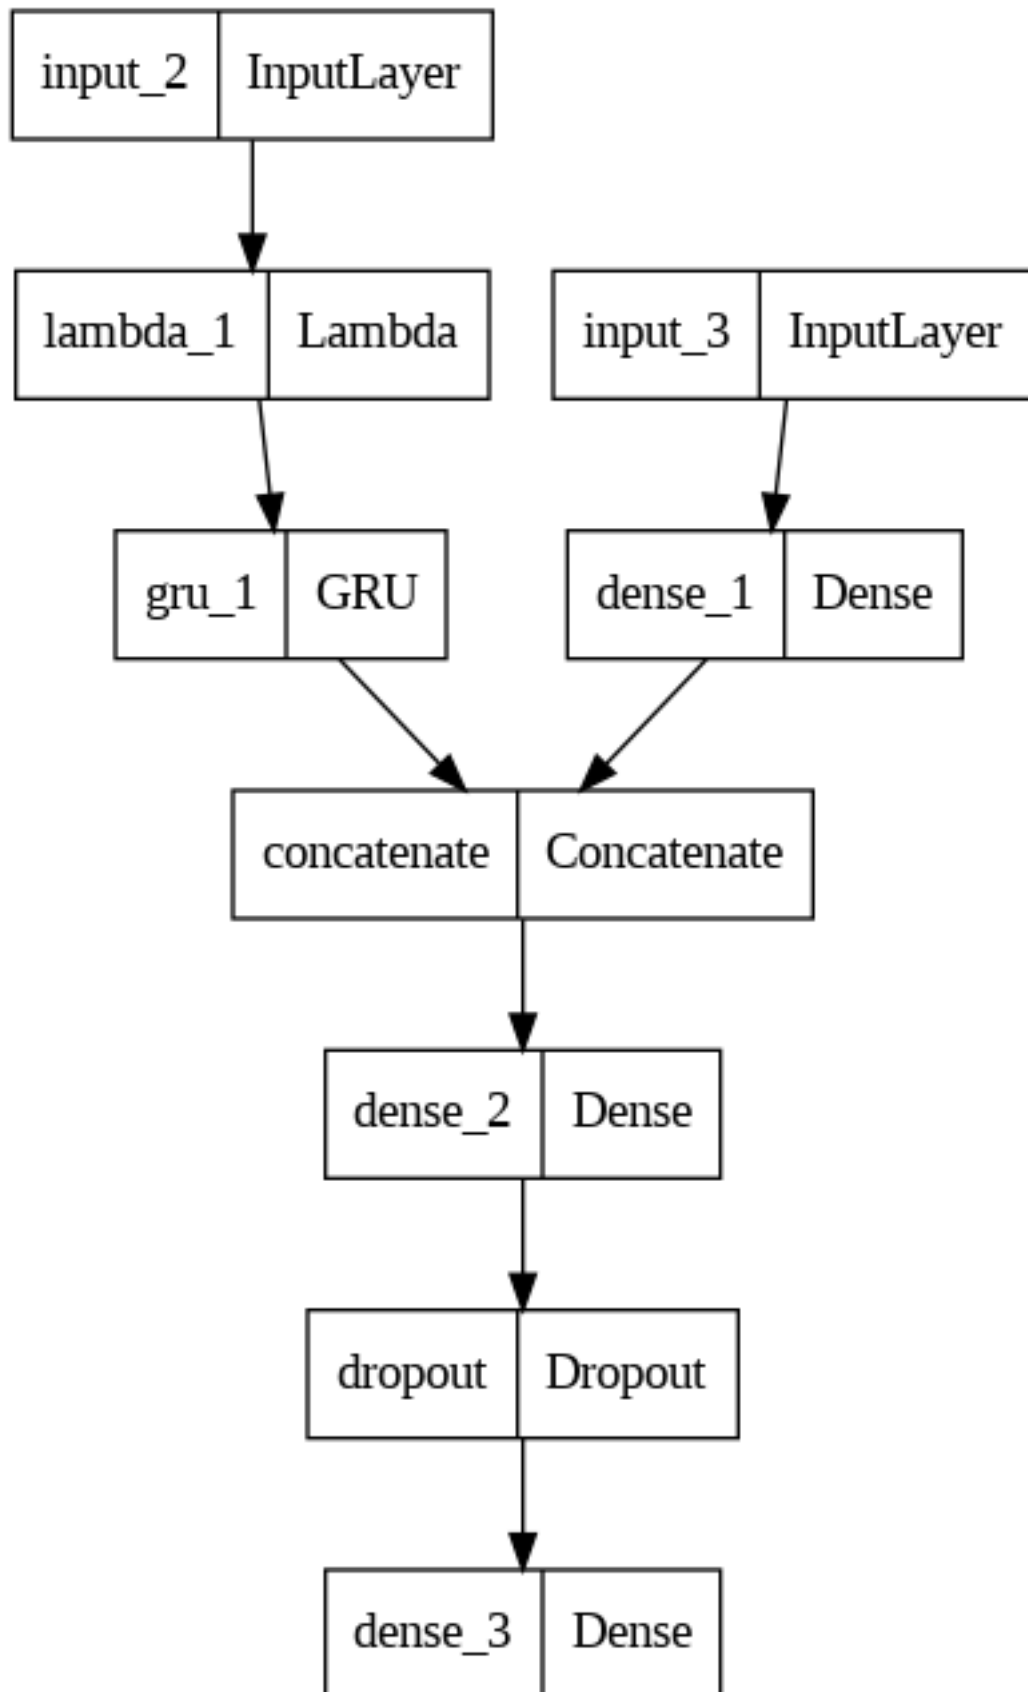
```
Total params: 25889 (101.13 KB)  
Trainable params: 25889 (101.13 KB)  
Non-trainable params: 0 (0.00 Byte)
```

```
-----
```

```
-----
```

```
[102]: tf.keras.utils.plot_model(model_3)
```

```
[102]:
```

```
[103]: from sklearn.preprocessing import LabelEncoder, OrdinalEncoder
X_cat_encoder = OrdinalEncoder(categories = [DAYS_OF_WEEK, MONTH_NAMES])
X_cat_encoded = X_cat_encoder.fit_transform(dataset2_cat_features)
X_cat_encoded, X_cat_encoder.categories_
```


```
[103]: (array([[2., 8.],
               [3., 8.],
               [4., 8.],
               ...,
               [1., 0.],
               [2., 0.],
               [3., 0.]]),
        [array(['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday',
                'Sunday'], dtype=object),
         array(['January', 'February', 'March', 'April', 'May', 'June', 'July',
                'August', 'September', 'October', 'November', 'December'],
                dtype=object)])
```

```
[104]: from tensorflow.data import Dataset

model3_history = model_3.fit(x=[dataset2_rv_history_features, X_cat_encoded],
    ↪ y=train_dataset2, epochs=5)
pd.DataFrame(model3_history.history).plot()
```

```
Epoch 1/5
61/61 [=====] - 4s 3ms/step - loss: 247.8772
Epoch 2/5
61/61 [=====] - 0s 4ms/step - loss: 95.5574
Epoch 3/5
61/61 [=====] - 0s 4ms/step - loss: 87.0351
Epoch 4/5
61/61 [=====] - 0s 4ms/step - loss: 81.4419
Epoch 5/5
61/61 [=====] - 0s 4ms/step - loss: 78.4773
```

```
[104]: <Axes: >
```



```
[106]: test_dataset2 = windowize_dataset(dataset[TEST_DATA_BOUNDARY_INDEX-WINDOW_SIZE:
↳].copy())
test_dataset2['Month.Name'] = pd.Series(test_dataset2.index,
↳index=test_dataset2.index)\
        .apply(lambda x: calendar.month_name[x.month])
test_dataset2 = test_dataset2.dropna()
test_dataset2
```

```
[106]:
```

	Row	Day	Day.Of.Week	Page.Loads	Unique.Visits	\
Date						
2020-01-16	1951	Thursday	5	3817	2951	
2020-01-17	1952	Friday	6	3175	2419	
2020-01-18	1953	Saturday	7	2336	1927	
2020-01-19	1954	Sunday	1	2597	2031	
2020-01-20	1955	Monday	2	3715	2948	
...	
2020-08-15	2163	Saturday	7	2221	1696	
2020-08-16	2164	Sunday	1	2724	2037	
2020-08-17	2165	Monday	2	3456	2638	
2020-08-18	2166	Tuesday	3	3581	2683	
2020-08-19	2167	Wednesday	4	2064	1564	

	First.Time.Visits	Returning.Visits	Returning.Visits[t-1]	\
Date				
2020-01-16	2510	441	394.0	
2020-01-17	2006	413	441.0	
2020-01-18	1681	246	413.0	
2020-01-19	1717	314	246.0	
2020-01-20	2505	443	314.0	
...	
2020-08-15	1373	323	386.0	
2020-08-16	1686	351	323.0	
2020-08-17	2181	457	351.0	
2020-08-18	2184	499	457.0	
2020-08-19	1297	267	499.0	

	Returning.Visits[t-2]	Returning.Visits[t-3]	Month.Name
Date			
2020-01-16	434.0	515.0	January
2020-01-17	394.0	434.0	January
2020-01-18	441.0	394.0	January
2020-01-19	413.0	441.0	January
2020-01-20	246.0	413.0	January
...
2020-08-15	458.0	427.0	August
2020-08-16	386.0	458.0	August
2020-08-17	323.0	386.0	August
2020-08-18	351.0	323.0	August
2020-08-19	457.0	351.0	August

[217 rows x 11 columns]

```
[107]: X_test_rv_history_input = test_dataset2[rv_cols]
X_test_rv_history_input
```

```
[107]:
```

	Returning.Visits[t-1]	Returning.Visits[t-2]	\
Date			
2020-01-16	394.0	434.0	
2020-01-17	441.0	394.0	
2020-01-18	413.0	441.0	
2020-01-19	246.0	413.0	
2020-01-20	314.0	246.0	
...	
2020-08-15	386.0	458.0	
2020-08-16	323.0	386.0	
2020-08-17	351.0	323.0	
2020-08-18	457.0	351.0	
2020-08-19	499.0	457.0	

Date	Returning.Visits[t-3]
2020-01-16	515.0
2020-01-17	434.0
2020-01-18	394.0
2020-01-19	441.0
2020-01-20	413.0
...	...
2020-08-15	427.0
2020-08-16	458.0
2020-08-17	386.0
2020-08-18	323.0
2020-08-19	351.0

[217 rows x 3 columns]

```
[108]: X_test_cat_input = test_dataset2[['Day', 'Month.Name']]
X_test_cat_input = X_cat_encoder.transform(X_test_cat_input)
X_test_cat_input.shape, X_test_cat_input[:5]
```

```
[108]: ((217, 2),
array([[3., 0.],
       [4., 0.],
       [5., 0.],
       [6., 0.],
       [0., 0.])))
```

```
[109]: model_3_preds = model_3.predict([X_test_rv_history_input, X_test_cat_input])
model_3_preds[:15]
```


7/7 [=====] - 1s 5ms/step

```
[109]: array([[305.5503 ],
       [449.51257],
       [340.03714],
       [218.20609],
       [506.13654],
       [457.20728],
       [470.06744],
       [389.448  ],
       [438.15656],
       [237.35915],
       [306.98065],
       [457.7482  ],
       [484.1681  ],
       [469.60846],
       [435.30145]], dtype=float32)
```

```
[110]: y_dataset = test_dataset2['Returning.Visits']
```

```
[111]: def evaluate_model_predictions(y_true, predictions, model_name):  
    metrics = evaluate_predictions(y_true, predictions)  
  
    MODEL_METRICS.loc[model_name] = metrics  
    plot_time_series(predictions.ravel(), start_index=1900)  
    return metrics  
  
evaluate_model_predictions(y_dataset, model_3_preds, 'model_3 (multi-input)')
```

```
[111]: {'mae': 66.37921923096829,  
      'mse': 7288.542556140803,  
      'rmse': 85.3729615050386,  
      'mape': 0.143966483112374}
```



```
[112]: MODEL_METRICS
```

```
[112]:
```

	mae	mse	rmse	mape
model_0	72.198157	8508.622120	92.242193	0.167139
model_1	88.477570	11614.890893	107.772403	0.194746
model_3 (multi-input)	66.379219	7288.542556	85.372962	0.143966

ensemble methods

```
[113]: def build_model_5(n_models, loss_fns):  
    models = []  
    for loss_fn in loss_fns:  
        print(f"Training {n_models} models for {loss_fn} loss...")  
        for i in range(n_models):  
            model = Sequential([  
                Input(shape=(WINDOW_SIZE,)),  
                Lambda(lambda x: tf.expand_dims(x, axis=1)),  
                GRU(128, activation='relu'),
```

```

        Dense(1, activation='linear')
    ])

    model.compile(loss=loss_fn, optimizer=tf.keras.optimizers.Adam())
    models.append(model)

return models

model_5 = build_model_5(n_models=5, loss_fns=['mae', 'mse', 'mape'])
model_5

```

Training 5 models for mae loss...
 Training 5 models for mse loss...
 Training 5 models for mape loss...

```

[113]: [<keras.src.engine.sequential.Sequential at 0x7be49a576f20>,
<keras.src.engine.sequential.Sequential at 0x7be49a559d20>,
<keras.src.engine.sequential.Sequential at 0x7be49a5757e0>,
<keras.src.engine.sequential.Sequential at 0x7be49a93fa00>,
<keras.src.engine.sequential.Sequential at 0x7be49a6f1db0>,
<keras.src.engine.sequential.Sequential at 0x7be499098220>,
<keras.src.engine.sequential.Sequential at 0x7be4990a8b50>,
<keras.src.engine.sequential.Sequential at 0x7be4990cd060>,
<keras.src.engine.sequential.Sequential at 0x7be49a5f3040>,
<keras.src.engine.sequential.Sequential at 0x7be4990c0850>,
<keras.src.engine.sequential.Sequential at 0x7be4990c2380>,
<keras.src.engine.sequential.Sequential at 0x7be4990ced10>,
<keras.src.engine.sequential.Sequential at 0x7be498f4a710>,
<keras.src.engine.sequential.Sequential at 0x7be498f4b8e0>,
<keras.src.engine.sequential.Sequential at 0x7be498f4b700>]

```

```

[114]: model_5[0].summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
lambda_2 (Lambda)	(None, 1, 3)	0
gru_2 (GRU)	(None, 128)	51072
dense_4 (Dense)	(None, 1)	129

Total params: 51201 (200.00 KB)
 Trainable params: 51201 (200.00 KB)

Non-trainable params: 0 (0.00 Byte)

```
[115]: for i, model in enumerate(model_5):  
        print(f"Training model {i+1} out of {len(model_5)} models")  
        model.fit(train_dataset, epochs=5, verbose=0)
```

```
Training model 1 out of 15 models  
Training model 2 out of 15 models  
Training model 3 out of 15 models  
Training model 4 out of 15 models  
Training model 5 out of 15 models  
Training model 6 out of 15 models  
Training model 7 out of 15 models  
Training model 8 out of 15 models  
Training model 9 out of 15 models  
Training model 10 out of 15 models  
Training model 11 out of 15 models  
Training model 12 out of 15 models  
Training model 13 out of 15 models  
Training model 14 out of 15 models  
Training model 15 out of 15 models
```

```
[116]: def ensemble_prediction(models):  
        predictions = []  
        for model in models:  
            pred = model.predict(test_dataset, verbose=0)  
            predictions.append(pred)  
  
        return np.array(predictions)  
  
model_5_all_preds = ensemble_prediction(model_5)  
model_5_all_preds.shape
```

WARNING:tensorflow:5 out of the last 14 calls to <function Model.make_predict_function.<locals>.predict_function at 0x7be495e8cd30> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

WARNING:tensorflow:6 out of the last 16 calls to <function Model.make_predict_function.<locals>.predict_function at 0x7be495e8d6c0> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2)

passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your `@tf.function` outside of the loop. For (2), `@tf.function` has `reduce_retracing=True` option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

[116]: (15, 217, 1)

[118]: `model_5_all_preds.shape`

[118]: (15, 217, 1)


```
[119]: def aggregate_ensemble_predictions(predictions):
        return tf.reduce_mean(predictions, axis=0).numpy()

model_5_preds = aggregate_ensemble_predictions(model_5_all_preds)
model_5_preds.shape
```

[119]: (217, 1)

[120]: `evaluate_model_predictions(y_true, model_5_preds, 'model_5 (ensemble)')`

```
[120]: {'mae': 73.04654872911867,
        'mse': 8556.55393296718,
        'rmse': 92.50164286631443,
        'mape': 0.16257963102721296}
```



[121]: `MODEL_METRICS`

	mae	mse	rmse	mape
model_0	72.198157	8508.622120	92.242193	0.167139
model_1	88.477570	11614.890893	107.772403	0.194746

model_3 (multi-input)	66.379219	7288.542556	85.372962	0.143966
model_5 (ensemble)	73.046549	8556.553933	92.501643	0.162580